

UDS Midterm Study Guide

Yabei Zeng

Table of contents

1	UDS Midterm Study Guide	2
1.1	Intro Chapter	2
1.1.1	Solving Problems with Data: No Order	2
1.1.2	Problem Refinement & Stakeholder Management	2
1.2	Proscriptive vs. Descriptive Questions	3
1.3	Exploratory Questions	3
1.3.1	Internal versus External Validity	4
1.4	Passive Prediction Questions	4
1.5	Causal Questions	5
1.5.1	Potential Outcomes Framework	6
1.5.2	Indicator Variables	7
1.5.3	Experiment	8
1.5.4	Power Calculations	10

1 UDS Midterm Study Guide

1.1 Intro Chapter

1.1.1 Solving Problems with Data: No Order

- **Step (1) Specifying the Problem**
- **Step (2) Solving Problems through Answering Questions**
 - Types of Questions:
 - * **Exploratory Questions:** Questions about general patterns in the data
 - Useful for understanding the problem space better and prioritizing subsequent efforts
 - E.g.: how many job applicants are you receiving when you post a job?
 - * **Passive Prediction Questions:** Questions about likely outcomes for individual observations or entities
 - Useful for targeting individuals for additional attention or interventions being considered
 - E.g.: given the symptoms of this patient and their test results, how likely are they to develop complications after surgery?
 - * **Causal Questions:** Questions about the effect of actions or interventions being considered
 - Useful for deciding on appropriate courses of action
 - E.g.: what is the effect of an action X on an outcome Y? If I do X, how will Y change?
 - Fundamental problem of causal inference: we can never perfectly know what the value of our outcome Y would be in both a world we do X and one where we don't do X

1.1.2 Problem Refinement & Stakeholder Management

- **Step (0) Recognize your role**
 - Helping your stakeholder better understand their problem is a core part of the job
 - Data science is about pairing domain expertise with computational methods and quantitative insights
- **Step (1) Don't assume your stakeholder knows what they need**
- **Step (2) Abstract the problem**
- **Step (3) Ask questions (Especially Quantitative Ones)**
 - Questions about Success

- Questions about the Problem
- **Step (4) Propose questions you might answer**
 - Make your questions specific and actionable
- **Step (5) Iterate**
 - Bring your work back to your stakeholder as often as possible

1.2 Proscriptive vs. Descriptive Questions

- Descriptive Questions: Questions about the **state of the world** and about objective reality (have right or wrong answers) include the following examples:
 - “What kinds of users are clicking our ads?”
 - “Do high-income and low-income countries emit similar amounts of carbon dioxide?”
- Proscriptive Questions: Questions about **how the world should be**, don’t have right or wrong answers because answers to proscriptive questions require evaluating the desirability of possible outcomes, include the following examples:
 - “Should higher income and low income countries be expected to meet the same carbon emission reduction standards?”
 - “Do high-income countries have a moral obligation to provide tuberculosis drugs to developing countries for free (or at cost)?”

1.3 Exploratory Questions

- Using Exploratory Questions: questions about elicit information, questions about broader patterns in the world, could be answered by simple summary statistics or plots, answering the exploratory questions makes you understand the contours of the problem you seek to solve, include the following examples:
 - What type of buildings (industrial, residential, commercial) consume the most power in the US?
 - In what region of the US are buildings consuming the most power?
 - Is there a region of the US where buildings are generating the most CO2?
 - Does the average energy use per building vary by region or building type?
 - In what season is most building energy consumed? Is more energy consumed by heating or AC needs, or do the two use similar amounts of power?

1.3.1 Internal versus External Validity

- Internal validity: how well you have analyzed the data you have
- External validity: how well you expect the answer you generated from that data to generalize to your stakeholder's context
- Internal validity concerns
 - Relate to your ability to properly characterize the causes of severe accidents in this data, accidents causes you generate from this data are meaning and faithful representations of the patterns in the data
 - Concerns over the accuracy with which things are measured
- External Validity Concerns
 - New features rolls out, whether there was anything exceptional about the data generating process
- The more control one has over a study context, the more likely one is to have good internal validity, but that control can often create an artificiality to limits external validity, Thus internal and external validity should not necessarily be thought of as things to try and simultaneously maximize at all costs, rather, they are best thought of as distinct features of analysis that should always be considered

1.4 Passive Prediction Questions

- Passive prediction questions: about **the future or potential** outcomes of individuals entities, including examples:
 - “How likely is patient X to experience a heart attack in the next two years?”
 - “How likely is it that Mortgage Holder Y will fail to make their mortgage payment next month?”
- Passive Prediction don't usually have “an answer”, need by considering the feasibility of developing a model to give individual-level answers to a Passive Prediction Question
- Differentiating between exploratory and passive prediction questions
 - Passive prediction questions focus on the values that get spit out of a model for each entity in the data, the only thing we care about in the passive prediction is **the quality of these predictions**
 - With exploratory question, our interest is in improving our understanding of the problem space
- Passive prediction and causal questions
 - Both questions are trying to predict future outcome

- The difference is that passive prediction are about make accurate prediction, not causal relationships that we can directly manipulate to shape outcomes (causal questions)
- Correlation does not necessarily imply causation
- **Internal validity and external validity** in answering passive prediction questions
 - Internal validity: measure of how well a model captures the meaningful variation in the data we already have
 - External validity: measure of how well we think that our model is likely to perform when faced with new data
 - Some of the factors that influence the External Validity to Passive Prediction Questions are the same as those that shape the External Validity of Exploratory Question
 - * Population represented in the data
 - * The time period in the questions
 - Two external validity concern (different from exploratory questions):
 - * Overfitting
 - * Adversarial users: attempt to subvert a statistical or machine learning model
- Not Using Black Box:
 - Problem with black box machine learning models: these models are used in high-stakes areas like **healthcare** and **criminal justice** but lack transparency, leading to severe consequences due to their inaccessibility and complexity
 - Explanation towards the black box could be problematic, lacking fidelity and can be misleading
 - Well-structured data with meaningful features can be equally well-predicted by simpler, interpretable models
- pyGAM
 - Generalized Additive Models(GAMs)
 - * $g(\mathbb{E}[y|X]) = \beta_0 + f_1(X_1) + f_2(X_2, X_3) + \dots + f_M(X_N)$
 - * X_1, X_2, \dots, X_N are independent variables, and y is the dependent variable
 - * $g()$ is the link function

1.5 Causal Questions

- Causal questions: designed to help us predict the consequences of our actions
- Stakeholder will ask a causal question when they want to know whether an action may be beneficial
- An answer to causal questions will involve designing a study that not only measures the effect of treatment but also knows any measured effect will generalize to the stakeholder's context.

- Internal and External Validity
 - Internal: the study itself in the context of studying
 - External: The context in which the stakeholder wishes to generalize the result
- Answering causal questions:
 - An action X on an outcome Y
 - **Fundamental Problem of Causal Inference:** directly measure the causal effect of X on Y for a given entity
 - **Randomized experiment:** most familiar tool for answering causal questions
 - * Randomized Control Trials (RCTs) or A/B Tests

1.5.1 Potential Outcomes Framework

- T is the treatment, a binary treatment, meaning it can only take on two values : $\{0,1\}$
- Y_i is the potential outcome variable for i individual, Y_i^1 is the treated potential outcome variable for i individual
- D is the observed receipt of either treatment or control, takes two values $\{0, 1\}$
- Given value of D, we will only be able to observe
 - Y_i^0 if $D_i = 0$ or Y_i^1 if $D_i = 1$
- Defining the causal effect:
 - $\delta_i = Y_i^1 - Y_i^0$
 - The difference between the individual potential outcome
 - This quantity is not able to observe
 - For all individuals, the value is call ATE
 - * $ATE = \frac{1}{N} \sum_{i \in 1,2,3...N} \delta_i$
 - * $ATE = \frac{1}{N} \sum_{i \in 1,2,3...N} Y_i^1 - Y_i^0$
 - * If we assume our data is a random sample from the population, then we can write the ATE in the following form
 - $ATE = E(\delta_i)$
 - $ATE = E(Y_i^1 - Y_i^0)$
 - $ATE = E(Y_i^1) - E(Y_i^0)$
 - * The binary treatment
 - $E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 0)$
 - This is a exact quantity you get from a linear regression of Y on D
- Correlation and Causation
 - $\hat{ATE} = E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 0) + E(Y_i^0|D_i = 1) - E(Y_i^0|D_i = 1)$
 - because $E(Y_i^0|D_i = 1) - E(Y_i^0|D_i = 1) = 0$, therefore $\hat{ATE} = E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 0)$

- We can also write \hat{ATE} as
 - * $\hat{ATE} = E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1) + E(Y_i^0|D_i = 1) - E(Y_i^0|D_i = 0)$
 - * where $E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1)$ is ATT
 - * and $E(Y_i^0|D_i = 1) - E(Y_i^0|D_i = 0)$ is the baseline difference
- Average Treatment Effect on the Treated (ATT)
 - $E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1)$
 - $ATE = E(Y_i^1) - E(Y_i^0)$
 - $= \lambda (E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1)) + (1 - \lambda) (E(Y_i^1|D_i = 0) - E(Y_i^0|D_i = 0))$
 - where $E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1)$ is the ATT (Average treatment effect on the treated)
 - where $E(Y_i^1|D_i = 0) - E(Y_i^0|D_i = 0)$ is the average treatment effect on the untreated
 - $ATT = ATE$ when $E(Y_i^1|D_i = 1) - E(Y_i^0|D_i = 1) = E(Y_i^1|D_i = 0) - E(Y_i^0|D_i = 0)$, meaning no differential treatment effects
 - neither $E(Y_i^0|D_i = 1)$ nor $E(Y_i^0|D_i = 0)$ are observable

WHAT is SUTVA?

- The Stable Unit Treatment Value Assumption (SUTVA) is a fundamental principle in causal inference that assumes **the treatment given to one unit does not affect the outcomes of other units**. This assumption comprises three main ideas:
 - **homogeneity of treatment effects**: each unit receives the same treatment effect
 - **No spillovers or externalities**: The treatment of one unit does not influence the outcomes of another unit.
 - **No general equilibrium effects**: The assumption that the treatment does not lead to broader systemic changes that could affect the outcomes.

1.5.2 Indicator Variables

Indicator variables: sometimes referred to as dummy variables, takes values of 0 and 1, used to indicate whether a given observation belongs to a discrete category in a way that can be used in statistical models.

- the coefficient on an indicator variable is an estimate of the average **DIFFERENCE** in the dependent variable for the group identified by the indicator variable
- the **REFERENCE GROUP**, is the set of observations for which the indicator variable is always zero.
- the coefficient on an indicator variable is an estimate of a **DIFFERENCE** with respect to a **REFERENCE GROUP**
- For categorical variables with more than 2 categories, use **one-hot encoding**

1.5.3 Experiment

A/B testing

A/B testing: two versions, and see which one performs better, includes the following steps:

- Identify the goal
- Create variants: two versions to be tested
- Split the audience: one group sees version A, and the other sees version B
- Collect the data
- Analyze results
- Make decisions

Twyman's Law: interesting or different results are often mistaken and emphasizes the need for skepticism and thorough verification of experiment data

- implications in Data Analysis:
 - error checking
 - skepticism
 - verification
- application in research and experimentation
 - design robustness
 - reproducibility
 - result interpretation

Overall Evaluation Criteria (OEC)

OEC: metrics used to measure the success of an experiment and determine whether the tested changes have achieved the desired outcome

- **Selection of OEC:** The right OEC involves the followings
 - considering the goals of the experiment
 - the expected impact of the changes
 - how success will be measured
- **Balance between short-term and long-term goals**
 - the balance should consider how changes might affect user behavior or business metrics over time

A/A Testing

Purpose/Definition: running two identical versions of a product or feature against each other to check the testing infrastructure's accuracy and to detect any biases or errors before conducting actual A/B tests

Through AA testing, a baseline for normal performance variation without experimental changes is established. The baseline is used for comparison in subsequent A/B tests to measure the impact of changes accurately

Internal Validity

Internal Validity: an experiment accurately establishes a causal relationship between the independent and dependent variables, ensuring the results are due to the manipulated variables and not other factors

Threats to internal validity

- selection bias
- history effect (external events affecting outcomes)
- maturation (participants changing over time)
- instrumentation changes (differences in measurement tools or procedures)

Strategies for improving internal validity

- random assignment to control and treatment groups
- controlling external variables
- using blinding methods
- ensuring consistent measurement techniques throughout the experiment

External Validity

external validity: how well the results of an experiment can be generalized beyond the specific conditions of the study

Factors affecting external validity

- sample characteristics
- experimental setting
- the context of which the experiment is conducted

Threats to external validity

- selection biases
- situational factors
- interaction effects between experimental treatments and participants

Strategies for improving external validity

- using representative samples
- ensure the experimental setting closely mirrors the real-world context
- replicating studies in different settings

1.5.4 Power Calculations

- **Statistical Power and Type II Errors:** defines statistical power as the probability of correctly rejecting a false null hypothesis and explains Type II errors (false negatives)
- **calculating sample size:** discusses how to compute the sample size needed to achieve desired power and control type I (false positives) and Type II error rates
- **Errors in Hypothesis Testing:**
 - True Positive (Correct Rejection of Null Hypothesis): The test correctly rejects the null hypothesis when it is false. This means that there is a real effect, and the test has successfully detected it.
 - True Negative (Correct Acceptance of Null Hypothesis): The test correctly fails to reject the null hypothesis when it is true. This means that there is no real effect, and the test correctly identifies that there is no significant difference or effect.
 - Type I Error (False Positive): The test incorrectly rejects the null hypothesis when it is actually true. This means the test indicates there is an effect when there isn't one, leading to a false alarm.
 - Type II Error (False Negative): The test incorrectly fails to reject the null hypothesis when it is actually false. This means the test fails to detect a real effect, missing the discovery of an actual difference or effect.
- **Type II Error Rate (β):** the probability of not detecting a true effect when it exists, decreases with larger sample sizes or greater effect sizes
- **Trade-offs in Testings:** the inverse relationship between Type I and Type II errors and the necessity of balancing sensitivity (power) against the risk of false positives.
- **Effect Size:** statistical power depends on the specified effect size, larger effects are easier to detect
- **Sample Size:** influenced by desired power, significance level (Type I error rate), and the minimum effect size of interest
- **Minimum Effect of Interest (MEI):** drive the design of an experiment based on the smallest practical effect size that would be meaningful to detect